

```

import gdown
file_id = '1rmJoBX-Hrm6MtPtjnVQ41LUXqytyARPI'
url = f'https://drive.google.com/uc?id={file_id}'
output = 'Major_5_District.csv'
gdown.download(url, output, quiet=False)
import pandas as pd
df = pd.read_csv('Major_5_District.csv')

→ Downloading...
From: https://drive.google.com/uc?id=1rmJoBX-Hrm6MtPtjnVQ41LUXqytyARPI
To: /content/Major_5_District.csv
100%|██████████| 57.8k/57.8k [00:00<00:00, 8.62MB/s]

```

▼ Gradient Boosting

df.columns

```

→ Index(['Year', 'Dist Name', 'Area', 'Production', 'Yield', 'Irrigated Area',
       'Annual Rainfall', 'NITROGEN CONSUMPTION (tons)',
       'NITROGEN SHARE IN NPK (Percent)', 'NITROGEN PER HA OF NCA (Kg per ha)',
       'NITROGEN PER HA OF GCA (Kg per ha)', 'PHOSPHATE CONSUMPTION (tons)',
       'PHOSPHATE SHARE IN NPK (Percent)',
       'PHOSPHATE PER HA OF NCA (Kg per ha)',
       'PHOSPHATE PER HA OF GCA (Kg per ha)', 'POTASH CONSUMPTION (tons)',
       'POTASH SHARE IN NPK (Percent)', 'POTASH PER HA OF NCA (Kg per ha)',
       'POTASH PER HA OF GCA (Kg per ha)', 'TOTAL CONSUMPTION (tons)',
       'TOTAL PER HA OF NCA (Kg per ha)', 'TOTAL AREA (1000 ha)',
       'FOREST AREA (1000 ha)', 'BARREN AND UNCULTIVABLE LAND AREA (1000 ha)',
       'LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha)',
       'CULTIVABLE WASTE LAND AREA (1000 ha)',
       'PERMANENT PASTURES AREA (1000 ha)', 'OTHER FALLOW AREA (1000 ha)',
       'CURRENT FALLOW AREA (1000 ha)', 'NET CROPPED AREA (1000 ha)',
       'GROSS CROPPED AREA (1000 ha)', 'CROPPING INTENSITY (Percent)',
       'Min Temp (Centigrade)', 'Max Temp (Centigrade)', 'Precipitation (mm)',
       'Evapotranspiration (mm)'],
      dtype='object')

```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

# Load the dataset
file_path = '/content/Major_5_District.csv'
data = pd.read_csv(file_path)
data = data.drop(columns=['TOTAL AREA (1000 ha)'])

# List of districts to forecast
districts = ['Akola', 'Kolhapur', 'Pune', 'Ratnagiri', 'Wardha'] # Replace with actual district
target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude 'Year' and 'Dist Name'

# Define a function to calculate additional metrics
def calculate_metrics(y_true, y_pred):
    r2 = r2_score(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_true, y_pred)
    mape = np.mean(np.abs((y_true - y_pred) / y_true)) * 100
    return {'R2': r2, 'MSE': mse, 'RMSE': rmse, 'MAE': mae, 'MAPE': mape}

# Dictionary to store results
results = {}

```

```

# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]
    results[district] = {}

    # Loop through each target column
    for target_column in target_columns:
        print(f"  Forecasting: {target_column}")

        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column])
        y = district_data[target_column]

        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Gradient Boosting Model
        model = GradientBoostingRegressor(random_state=42)
        model.fit(X_train, y_train)

        # Predictions
        y_pred = model.predict(X_test)

        # Evaluate the model
        metrics = calculate_metrics(y_test, y_pred)
        results[district][target_column] = metrics

# Print results
for district, district_metrics in results.items():
    print(f"\nResults for {district}:")
    for column, metrics in district_metrics.items():
        print(f"  {column}:")
        for metric, value in metrics.items():
            print(f"    {metric}: {value:.4f}")

```

→ Processing District: Akola

Forecasting: Area

Forecasting: Production

Forecasting: Yield

Forecasting: Irrigated Area

Forecasting: Annual Rainfall

Forecasting: NITROGEN CONSUMPTION (tons)

Forecasting: NITROGEN SHARE IN NPK (Percent)

Forecasting: NITROGEN PER HA OF NCA (Kg per ha)

Forecasting: NITROGEN PER HA OF GCA (Kg per ha)

Forecasting: PHOSPHATE CONSUMPTION (tons)

Forecasting: PHOSPHATE SHARE IN NPK (Percent)

Forecasting: PHOSPHATE PER HA OF NCA (Kg per ha)

Forecasting: PHOSPHATE PER HA OF GCA (Kg per ha)

Forecasting: POTASH CONSUMPTION (tons)

Forecasting: POTASH SHARE IN NPK (Percent)

Forecasting: POTASH PER HA OF NCA (Kg per ha)

Forecasting: POTASH PER HA OF GCA (Kg per ha)

Forecasting: TOTAL CONSUMPTION (tons)

Forecasting: TOTAL PER HA OF NCA (Kg per ha)

Forecasting: FOREST AREA (1000 ha)

Forecasting: BARREN AND UNCULTIVABLE LAND AREA (1000 ha)

Forecasting: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha)

Forecasting: CULTIVABLE WASTE LAND AREA (1000 ha)

Forecasting: PERMANENT PASTURES AREA (1000 ha)

Forecasting: OTHER FALLOW AREA (1000 ha)

Forecasting: CURRENT FALLOW AREA (1000 ha)

Forecasting: NET CROPPED AREA (1000 ha)

Forecasting: GROSS CROPPED AREA (1000 ha)

```

Forecasting: CROPPING INTENSITY (Percent)
Forecasting: Min Temp (Centigrade)
Forecasting: Max Temp (Centigrade)
Forecasting: Precipitation (mm)
Forecasting: Evapotranspiration (mm)
Processing District: Kolhapur
Forecasting: Area
Forecasting: Production
Forecasting: Yield
Forecasting: Irrigated Area
Forecasting: Annual Rainfall
Forecasting: NITROGEN CONSUMPTION (tons)
Forecasting: NITROGEN SHARE IN NPK (Percent)
Forecasting: NITROGEN PER HA OF NCA (Kg per ha)
Forecasting: NITROGEN PER HA OF GCA (Kg per ha)
Forecasting: PHOSPHATE CONSUMPTION (tons)
Forecasting: PHOSPHATE SHARE IN NPK (Percent)
Forecasting: PHOSPHATE PER HA OF NCA (Kg per ha)
Forecasting: PHOSPHATE PER HA OF GCA (Kg per ha)
Forecasting: POTASH CONSUMPTION (tons)
Forecasting: POTASH SHARE IN NPK (Percent)
Forecasting: POTASH PER HA OF NCA (Kg per ha)
Forecasting: POTASH PER HA OF GCA (Kg per ha)
Forecasting: TOTAL CONSUMPTION (tons)
Forecasting: TOTAL PER HA OF NCA (Kg per ha)
Forecasting: FOREST AREA (1000 ha)
Forecasting: BARREN AND UNCULTIVABLE LAND AREA (1000 ha)
Forecasting: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha)
Forecasting: CIIITTVARI F WASTF I AND ARFA (1000 ha)

# Dictionary to store all forecasted values
forecasted_values = []

# Iterate again over each district to generate forecasted data
for district in districts:
    district_data = data[data['Dist Name'] == district]

    for target_column in target_columns:
        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column])
        y = district_data[target_column]

        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Gradient Boosting Model
        model = GradientBoostingRegressor(random_state=42)
        model.fit(X_train, y_train)

        # Predictions for the test set
        y_pred = model.predict(X_test)

        # Store results
        for i in range(len(X_test)):
            forecasted_values.append({
                'District': district,
                'Target Column': target_column,
                'Actual Value': y_test.iloc[i],
                'Predicted Value': y_pred[i]
            })

# Convert forecasted values to DataFrame
forecasted_df = pd.DataFrame(forecasted_values)

# Save to CSV
forecasted_csv_path = '/content/forecasted_values.csv'
forecasted_df.to_csv(forecasted_csv_path, index=False)

```

```
print(f"Forecasted values saved to {forecasted_csv_path}")

➡️ Forecasted values saved to /content/forecasted_values.csv

# Prepare data for performance metrics
metrics_list = []

for district, district_metrics in results.items():
    for column, metrics in district_metrics.items():
        metrics_row = {'District': district, 'Target Column': column}
        metrics_row.update(metrics) # Add all metrics
        metrics_list.append(metrics_row)

# Convert metrics list to DataFrame
metrics_df = pd.DataFrame(metrics_list)

# Save to CSV
metrics_csv_path = '/content/performance_metrics.csv'
metrics_df.to_csv(metrics_csv_path, index=False)

print(f"Performance metrics saved to {metrics_csv_path}")
```

```
➡️ Performance metrics saved to /content/performance_metrics.csv

# Initialize a list to store aggregate metrics
aggregate_metrics = []

# Iterate through each district
for district in results.keys():
    district_metrics = results[district]

    # Initialize accumulators for each metric
    r2_sum = 0
    mse_sum = 0
    rmse_sum = 0
    mae_sum = 0
    mape_sum = 0
    count = 0

    # Aggregate metrics across all target columns
    for target_column, metrics in district_metrics.items():
        r2_sum += metrics['R2']
        mse_sum += metrics['MSE']
        rmse_sum += metrics['RMSE']
        mae_sum += metrics['MAE']
        mape_sum += metrics['MAPE']
        count += 1

    # Calculate averages
    aggregate_metrics.append({
        'District': district,
        'Avg R²': r2_sum / count,
        'Avg MSE': mse_sum / count,
        'Avg RMSE': rmse_sum / count,
        'Avg MAE': mae_sum / count,
```

```

        'Avg MAPE': mape_sum / count
    })

# Convert aggregate metrics to a DataFrame
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)

# Save to CSV if needed
aggregate_metrics_csv_path = '/content/aggregate_performance_metrics.csv'
aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)

print(f"Aggregate performance metrics saved to {aggregate_metrics_csv_path}")

# Display the DataFrame
aggregate_metrics_df

```

Aggregate performance metrics saved to /content/aggregate_performance_metrics.csv

	District	Avg R ²	Avg MSE	Avg RMSE	Avg MAE	Avg MAPE
0	Akola	0.547307	1.004943e+06	366.135866	268.107496	11.079484
1	Kolhapur	0.679010	2.239698e+06	509.385907	384.459961	7.548004
2	Pune	0.745560	1.970016e+06	549.471303	394.446955	10.711596
3	Ratnagiri	0.515195	9.109478e+05	280.685871	208.320852	8.678274
4	Wardha	0.690796	8.265315e+05	335.163144	260.200483	14.867005

▼ Forecasting

```

import pandas as pd
import numpy as np
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split

# Load the dataset
file_path = '/content/Major_5_District.csv'
data = pd.read_csv(file_path)
data = data.drop(columns=['TOTAL AREA (1000 ha)'])

# Years to forecast
future_years = list(range(2018, 2026)) + [2030, 2035, 2040, 2045]

# List of districts to forecast
districts = ['Akola', 'Kolhapur', 'Pune', 'Ratnagiri', 'Wardha'] # Replace with actual district
target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude 'Year' and 'Dist Name'

# Dictionary to store forecasted results
forecasted_results = []

# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]

    for target_column in target_columns:
        print(f"Forecasting: {target_column}")

        # Prepare features (X) and target (y)

```

```

X = district_data.drop(columns=[ 'Year', 'Dist Name', target_column])
y = district_data[target_column]

# Train-test split (train on all historical data for forecasting)
X_train = X
y_train = y

# Gradient Boosting Model
model = GradientBoostingRegressor(random_state=42)
model.fit(X_train, y_train)

# Create future data for predictions
future_data = district_data.iloc[:len(future_years)].copy()
future_data['Year'] = future_years
future_X = future_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')

# Predict future values
future_predictions = model.predict(future_X)

# Store results
for year, prediction in zip(future_years, future_predictions):
    forecasted_results.append({
        'District': district,
        'Year': year,
        'Target Column': target_column,
        'Forecasted Value': prediction
    })

# Convert forecasted results to DataFrame
forecasted_df = pd.DataFrame(forecasted_results)

# Save to CSV
forecasted_csv_path = '/content/forecasted_future_values.csv'
forecasted_df.to_csv(forecasted_csv_path, index=False)

print(f"Forecasted values saved to {forecasted_csv_path}")

```

→ Processing District: Akola
 Forecasting: Area
 Forecasting: Production
 Forecasting: Yield
 Forecasting: Irrigated Area
 Forecasting: Annual Rainfall
 Forecasting: NITROGEN CONSUMPTION (tons)
 Forecasting: NITROGEN SHARE IN NPK (Percent)
 Forecasting: NITROGEN PER HA OF NCA (Kg per ha)
 Forecasting: NITROGEN PER HA OF GCA (Kg per ha)
 Forecasting: PHOSPHATE CONSUMPTION (tons)
 Forecasting: PHOSPHATE SHARE IN NPK (Percent)
 Forecasting: PHOSPHATE PER HA OF NCA (Kg per ha)
 Forecasting: PHOSPHATE PER HA OF GCA (Kg per ha)
 Forecasting: POTASH CONSUMPTION (tons)
 Forecasting: POTASH SHARE IN NPK (Percent)
 Forecasting: POTASH PER HA OF NCA (Kg per ha)
 Forecasting: POTASH PER HA OF GCA (Kg per ha)
 Forecasting: TOTAL CONSUMPTION (tons)
 Forecasting: TOTAL PER HA OF NCA (Kg per ha)
 Forecasting: FOREST AREA (1000 ha)
 Forecasting: BARREN AND UNCULTIVABLE LAND AREA (1000 ha)
 Forecasting: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha)
 Forecasting: CULTIVABLE WASTE LAND AREA (1000 ha)
 Forecasting: PERMANENT PASTURES AREA (1000 ha)
 Forecasting: OTHER FALLOW AREA (1000 ha)
 Forecasting: CURRENT FALLOW AREA (1000 ha)
 Forecasting: NET CROPPED AREA (1000 ha)
 Forecasting: GROSS CROPPED AREA (1000 ha)

```
Forecasting: CROPPING INTENSITY (Percent)
Forecasting: Min Temp (Centigrade)
Forecasting: Max Temp (Centigrade)
Forecasting: Precipitation (mm)
Forecasting: Evapotranspiration (mm)
Processing District: Kolhapur
Forecasting: Area
Forecasting: Production
Forecasting: Yield
Forecasting: Irrigated Area
Forecasting: Annual Rainfall
Forecasting: NITROGEN CONSUMPTION (tons)
Forecasting: NITROGEN SHARE IN NPK (Percent)
Forecasting: NITROGEN PER HA OF NCA (Kg per ha)
Forecasting: NITROGEN PER HA OF GCA (Kg per ha)
Forecasting: PHOSPHATE CONSUMPTION (tons)
Forecasting: PHOSPHATE SHARE IN NPK (Percent)
Forecasting: PHOSPHATE PER HA OF NCA (Kg per ha)
Forecasting: PHOSPHATE PER HA OF GCA (Kg per ha)
Forecasting: POTASH CONSUMPTION (tons)
Forecasting: POTASH SHARE IN NPK (Percent)
Forecasting: POTASH PER HA OF NCA (Kg per ha)
Forecasting: POTASH PER HA OF GCA (Kg per ha)
Forecasting: TOTAL CONSUMPTION (tons)
Forecasting: TOTAL PER HA OF NCA (Kg per ha)
Forecasting: FOREST AREA (1000 ha)
Forecasting: BARREN AND UNCULTIVABLE LAND AREA (1000 ha)
Forecasting: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha)
Forecasting: CULTIVATED LAND AREA (1000 ha)
```

Plot

```
import matplotlib.pyplot as plt
import seaborn as sns

# Ensure the plots use seaborn style
sns.set(style="whitegrid")

# Iterate over districts to create individual plots
districts = forecasted_df['District'].unique()

for district in districts:
    # Filter data for the specific district
    district_data = forecasted_df[forecasted_df['District'] == district]

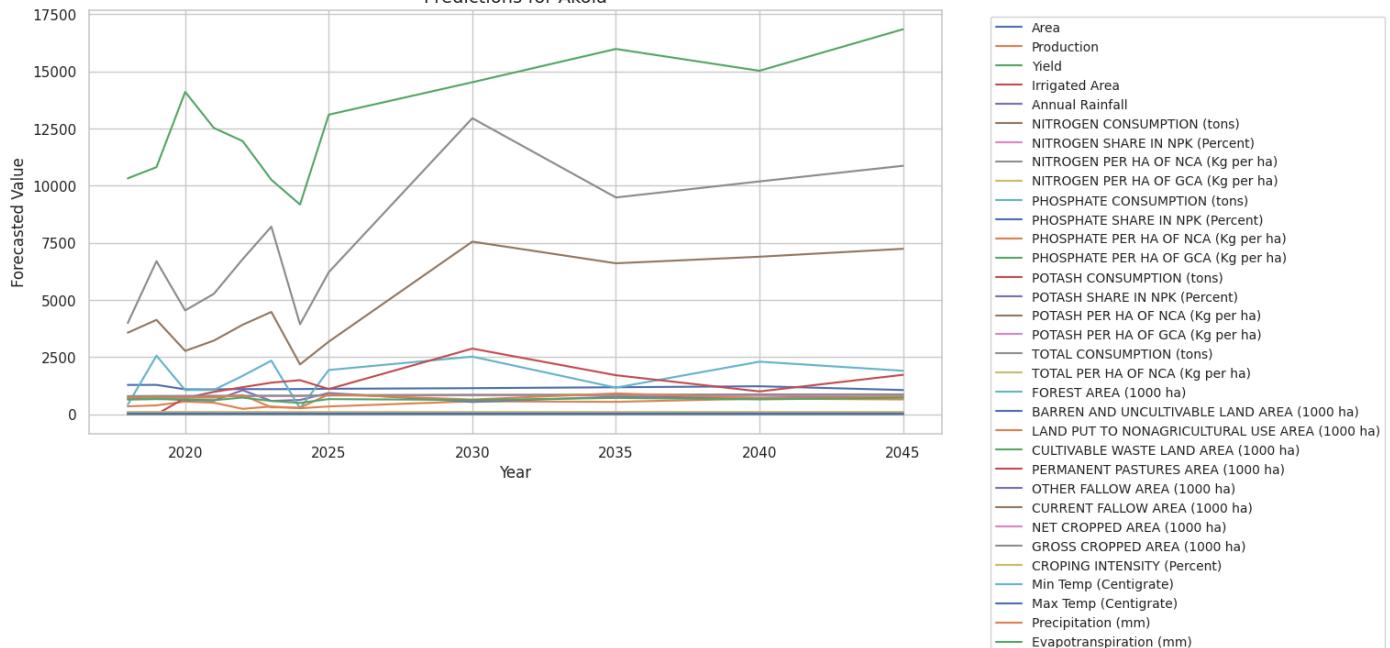
    # Create a new plot for the district
    plt.figure(figsize=(12, 6))

    # Plot each target column
    target_columns = district_data['Target Column'].unique()
    for target in target_columns:
        target_data = district_data[district_data['Target Column'] == target]
        plt.plot(target_data['Year'], target_data['Forecasted Value'], label=target)

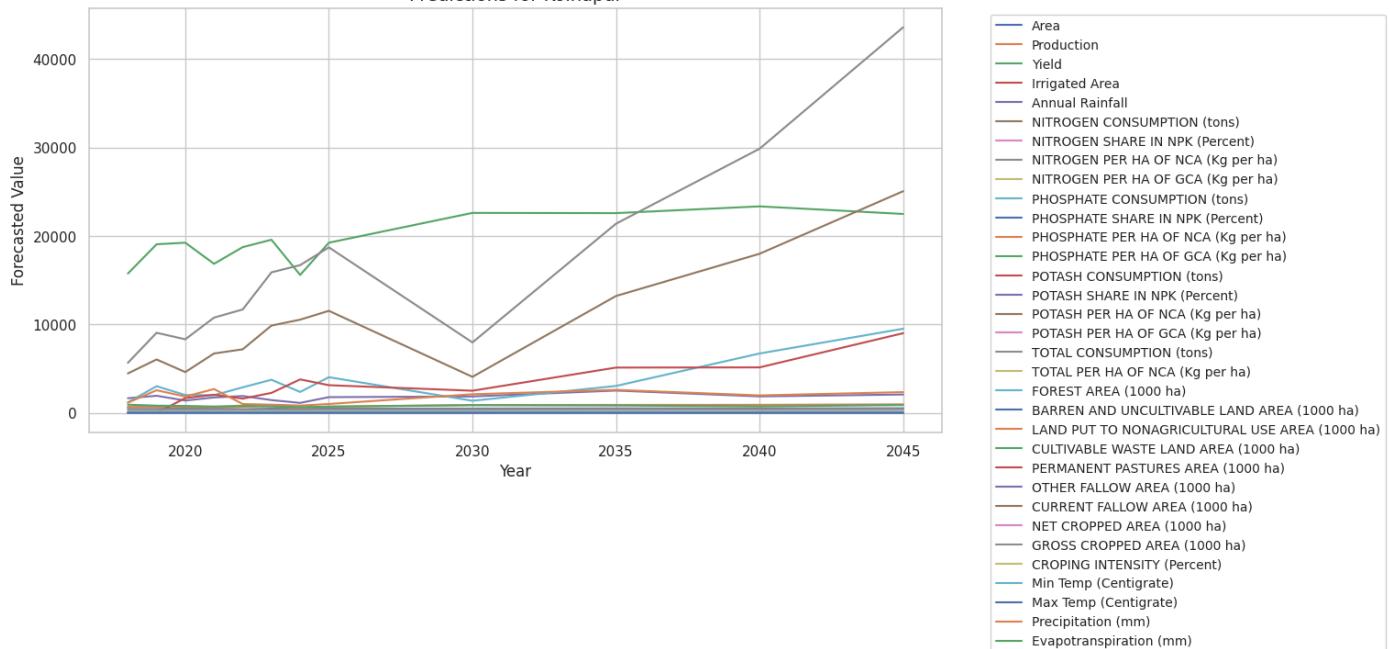
    # Add title, labels, and legend
    plt.title(f"Predictions for {district}", fontsize=14)
    plt.xlabel("Year", fontsize=12)
    plt.ylabel("Forecasted Value", fontsize=12)
    plt.legend(loc="upper left", bbox_to_anchor=(1.05, 1), fontsize=10) # Legend outside the plot
    plt.show()
```



Predictions for Akola



Predictions for Kolhapur



Finetuning

Predictions for Pune

```

from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import mean_squared_error
import pandas as pd
import numpy as np

# Define districts and target columns
districts = ['Akola', 'Kolhapur', 'Pune', 'Ratnagiri', 'Wardha']
target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude non-target columns

# Dictionary to store results
fine_tuning_results = []

# Loop through each district

```

```

for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]

    # Loop through each target column
    for target_column in target_columns:
        print(f"  Fine-tuning for Target Column: {target_column}")

        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')
        y = district_data[target_column]

        # Skip if there is insufficient data
        if len(y) < 5:
            print(f"    Skipping {target_column} due to insufficient data.")
            continue

        # Split data into training and testing sets
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Define the hyperparameter grid
        param_grid = {
            'n_estimators': [100, 200, 300],
            'learning_rate': [0.01, 0.05, 0.1],
            'max_depth': [3, 4, 5],
            'min_samples_split': [2, 5, 10],
            'min_samples_leaf': [1, 3, 5],
            'subsample': [0.8, 1.0]
        }

        # Instantiate the model
        gbr = GradientBoostingRegressor(random_state=42)

        # Use GridSearchCV to find the best parameters
        grid_search = GridSearchCV(estimator=gbr, param_grid=param_grid, cv=3, scoring='neg_mean_squared_error')
        grid_search.fit(X_train, y_train)

        # Best parameters and evaluation
        best_params = grid_search.best_params_
        best_model = grid_search.best_estimator_
        y_pred = best_model.predict(X_test)

        # Calculate metrics
        mse = mean_squared_error(y_test, y_pred)
        rmse = np.sqrt(mse)
        r2 = best_model.score(X_test, y_test)

        print(f"    RMSE: {rmse:.4f}, R2: {r2:.4f}")

        # Store results
        fine_tuning_results.append({
            'District': district,
            'Target Column': target_column,
            'Best Parameters': best_params,
            'RMSE': rmse,
            'R2

```

```

# Convert results to a DataFrame
results_df = pd.DataFrame(fine_tuning_results)

# Save results to a CSV
results_csv_path = '/content/fine_tuning_results.csv'
results_df.to_csv(results_csv_path, index=False)

print(f"Fine-tuning results saved to {results_csv_path}")

```

```

→ Processing District: Akola
  Fine-tuning for Target Column: Area
    RMSE: 158.4976, R2: 0.4095
  Fine-tuning for Target Column: Production
-----
KeyboardInterrupt                                     Traceback (most recent call last)
<ipython-input-19-6bc7b730baab> in <cell line: 15>()
    48     # Use GridSearchCV to find the best parameters
    49     grid_search = GridSearchCV(estimator=gbr, param_grid=param_grid, cv=3, scoring='neg_mean_squared_error', n_jobs=-1,
verbose=0)
--> 50     grid_search.fit(X_train, y_train)
    51
    52     # Best parameters and evaluation

_____
↓ 7 frames ↓
/usr/local/lib/python3.10/dist-packages/joblib/parallel.py in _retrieve(self)
1760         (self._jobs[0].get_status(
1761             timeout=self.timeout) == TASK_PENDING)):
-> 1762         time.sleep(0.01)
1763         continue
1764

KeyboardInterrupt:

```

▼ Hail Mary

```

from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd
import numpy as np
from joblib import Parallel, delayed

# Load the dataset
file_path = '/content/Major_5_District.csv'
data = pd.read_csv(file_path)
data = data.drop(columns=['TOTAL AREA (1000 ha)']) # Adjust as per actual column name

# Future years to forecast
future_years = list(range(2018, 2026)) + [2030, 2035, 2040, 2045]

# Districts to process
districts = ['Akola', 'Kolhapur', 'Pune', 'Ratnagiri', 'Wardha']
target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude non-target columns

# Hyperparameter tuning grid
param_distributions = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 5],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 3],
    'subsample': [0.8, 1.0]
}

```

```

}

# Function to process each district and target column
def process_district(district):
    results = {'fine_tuning': [], 'forecasting': []}
    district_data = data[data['Dist Name'] == district]

    for target_column in target_columns:
        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')
        y = district_data[target_column]

        # Skip columns with insufficient data
        if len(y) < 5:
            continue

        # Train on historical data
        X_train = X
        y_train = y

        # Randomized search for hyperparameter tuning
        gbr = GradientBoostingRegressor(random_state=42)
        random_search = RandomizedSearchCV(
            estimator=gbr,
            param_distributions=param_distributions,
            n_iter=10, # Reducing iterations for speed
            cv=2,
            scoring='neg_mean_squared_error',
            n_jobs=-1,
            verbose=0
        )
        random_search.fit(X_train, y_train)

        # Best parameters and model
        best_params = random_search.best_params_
        best_model = random_search.best_estimator_

        # Forecasting future years
        future_data = district_data.iloc[:len(future_years)].copy()
        future_data['Year'] = future_years
        future_X = future_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')
        future_predictions = best_model.predict(future_X)

        # Store fine-tuning results
        mse = mean_squared_error(y_train, best_model.predict(X_train))
        r2 = r2_score(y_train, best_model.predict(X_train))
        results['fine_tuning'].append({
            'District': district,
            'Target Column': target_column,
            'Best Parameters': best_params,
            'MSE': mse,
            'R2

```

```

        'Year': year,
        'Target Column': target_column,
        'Forecasted Value': prediction
    })

return results

# Run processing in parallel
all_results = Parallel(n_jobs=-1)(delayed(process_district)(district) for district in districts)

# Combine results
fine_tuning_results = [item for result in all_results for item in result['fine_tuning']]
forecasted_results = [item for result in all_results for item in result['forecasting']]

# Save fine-tuning metrics to CSV
metrics_df = pd.DataFrame(fine_tuning_results)
metrics_csv_path = '/content/fine_tuning_results.csv'
metrics_df.to_csv(metrics_csv_path, index=False)
print(f"Fine-tuning results saved to {metrics_csv_path}")

# Save forecasted results to CSV
forecasted_df = pd.DataFrame(forecasted_results)
forecasted_csv_path = '/content/forecasted_future_values.csv'
forecasted_df.to_csv(forecasted_csv_path, index=False)
print(f"Forecasted values saved to {forecasted_csv_path}")

```

→ Fine-tuning results saved to /content/fine_tuning_results.csv
Forecasted values saved to /content/forecasted_future_values.csv

```

# Initialize lists for aggregate metrics and forecast summaries
aggregate_metrics = []
forecast_summaries = []

# Aggregate fine-tuning results by district
for district in metrics_df['District'].unique():
    district_metrics = metrics_df[metrics_df['District'] == district]

    # Calculate average MSE and R2 for the district
    avg_mse = district_metrics['MSE'].mean()
    avg_r2 = district_metrics['R2'].mean()

    # Append aggregate metrics for the district
    aggregate_metrics.append({
        'District': district,
        'Avg MSE': avg_mse,
        'Avg R2

```

```

# Append forecast summary for the district and year
forecast_summaries.append({
    'District': district,
    'Year': year,
    'Avg Forecasted Value': avg_forecast
})

# Convert aggregate metrics and forecast summaries to DataFrames
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)
forecast_summary_df = pd.DataFrame(forecast_summaries)

# Save aggregated metrics to CSV
aggregate_metrics_csv_path = '/content/aggregate_fine_tuning_metrics.csv'
aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)
print(f"Aggregate fine-tuning metrics saved to {aggregate_metrics_csv_path}")

# Save forecast summaries to CSV
forecast_summary_csv_path = '/content/forecast_summary.csv'
forecast_summary_df.to_csv(forecast_summary_csv_path, index=False)
print(f"Forecast summaries saved to {forecast_summary_csv_path}")

# Display the DataFrames
print("\nAggregate Fine-Tuning Metrics:")
print(aggregate_metrics_df)

print("\nForecast Summary:")
print(forecast_summary_df)

```

→ Aggregate fine-tuning metrics saved to /content/aggregate_fine_tuning_metrics.csv
 Forecast summaries saved to /content/forecast_summary.csv

Aggregate Fine-Tuning Metrics:

	District	Avg MSE	Avg R ²
0	Akola	2254.957117	0.977871
1	Kolhapur	173.519715	0.955560
2	Pune	4612.143234	0.954814
3	Ratnagiri	136278.467375	0.950932
4	Wardha	297.810368	0.955835

Forecast Summary:

	District	Year	Avg Forecasted Value
0	Akola	2018	741.171287
1	Akola	2019	915.279504
2	Akola	2020	876.541802
3	Akola	2021	871.306690
4	Akola	2022	958.421581
5	Akola	2023	964.419282
6	Akola	2024	674.186669
7	Akola	2025	961.629556
8	Akola	2030	1396.317802
9	Akola	2035	1246.929210
10	Akola	2040	1264.442005
11	Akola	2045	1360.813980
12	Kolhapur	2018	1019.384957
13	Kolhapur	2019	1379.454188
14	Kolhapur	2020	1295.320752
15	Kolhapur	2021	1408.105150
16	Kolhapur	2022	1473.592296
17	Kolhapur	2023	1742.433403
18	Kolhapur	2024	1651.036059
19	Kolhapur	2025	1915.153772
20	Kolhapur	2030	1405.569863
21	Kolhapur	2035	2257.380928
22	Kolhapur	2040	2757.982042
23	Kolhapur	2045	3584.665112
24	Pune	2018	1072.613311
25	Pune	2019	1469.877096
26	Pune	2020	1175.535973
27	Pune	2021	1198.904664

```

28     Pune 2022      1443.719994
29     Pune 2023      1509.780069
30     Pune 2024      1490.125363
31     Pune 2025      1844.445292
32     Pune 2030      1786.620069
33     Pune 2035      1496.244709
34     Pune 2040      1880.540247
35     Pune 2045      1638.356534
36 Ratnagiri 2018      804.320929
37 Ratnagiri 2019      920.650142
38 Ratnagiri 2020      896.443478
39 Ratnagiri 2021      921.148232
40 Ratnagiri 2022      977.503991
41 Ratnagiri 2023      1013.426245
42 Ratnagiri 2024      852.681707
43 Ratnagiri 2025      1035.113936
44 Ratnagiri 2030      868.411059

# Initialize a list to store aggregate metrics
aggregate_metrics = []

# Get unique districts from the results
unique_districts = results_df['District'].unique()

# Iterate through each district
for district in unique_districts:
    # Filter results for the current district
    district_data = results_df[results_df['District'] == district]

    # Calculate average metrics for the district
    avg_rmse = district_data['RMSE'].mean()
    avg_r2 = district_data['R^2'].mean()

    # Append aggregate metrics for the district
    aggregate_metrics.append({
        'District': district,
        'Avg RMSE': avg_rmse,
        'Avg R^2': avg_r2
    })

# Convert aggregate metrics to a DataFrame
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)

# Save to CSV if needed
aggregate_metrics_csv_path = '/content/aggregate_performance_metrics.csv'
aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)

print(f"Aggregate performance metrics saved to {aggregate_metrics_csv_path}")

# Display the DataFrame
aggregate_metrics_df

# # Inside the process_district function, after finding the best hyperparameters
# random_search.fit(X_train, y_train)

# # Best parameters and model
# best_params = random_search.best_params_
# best_model = random_search.best_estimator_

# # Print the best hyperparameters for the current district and target column
# print(f"Best hyperparameters for {district} - {target_column}: {best_params}")

```

```
# Print the best hyperparameters for each district and target column
for result in fine_tuning_results:
    district = result['District']
    target_column = result['Target Column']
    best_params = result['Best Parameters']
    print(f"District: {district}, Target Column: {target_column}, Best Hyperparameters: {best_params}")
```

☒ District: Akola, Target Column: Area, Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_samples_split': 5, 'min_sa
 District: Akola, Target Column: Production, Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 5, 'mi
 District: Akola, Target Column: Yield, Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_samples_split': 5, 'min_s
 District: Akola, Target Column: Irrigated Area, Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_samples_split':
 District: Akola, Target Column: Annual Rainfall, Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'min_samples_split':
 District: Akola, Target Column: NITROGEN CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_sa
 District: Akola, Target Column: NITROGEN SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 300, 'min
 District: Akola, Target Column: NITROGEN PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'mi
 District: Akola, Target Column: NITROGEN PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'mi
 District: Akola, Target Column: PHOSPHATE CONSUMPTION (tons), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 300, 'min_sa
 District: Akola, Target Column: PHOSPHATE SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 300, 'mi
 District: Akola, Target Column: PHOSPHATE PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'mi
 District: Akola, Target Column: PHOSPHATE PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'mi
 District: Akola, Target Column: POTASH CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_sa
 District: Akola, Target Column: POTASH SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_s
 District: Akola, Target Column: POTASH PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'mi
 District: Akola, Target Column: POTASH PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100, 'mi
 District: Akola, Target Column: TOTAL CONSUMPTION (tons), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_sample
 District: Akola, Target Column: TOTAL PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min
 District: Akola, Target Column: FOREST AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_samples_s
 District: Akola, Target Column: BARREN AND UNCULTIVABLE LAND AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators
 District: Akola, Target Column: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators
 District: Akola, Target Column: CULTIVABLE WASTE LAND AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100,
 District: Akola, Target Column: PERMANENT PASTURES AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'm
 District: Akola, Target Column: OTHER FALLOW AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_sa
 District: Akola, Target Column: CURRENT FALLOW AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'min_s
 District: Akola, Target Column: NET CROPPED AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_samp
 District: Akola, Target Column: GROSS CROPPED AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_sa
 District: Akola, Target Column: CROPPING INTENSITY (Percent), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_sa
 District: Akola, Target Column: Min Temp (Centigrade), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100, 'min_samples_s
 District: Akola, Target Column: Max Temp (Centigrade), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'min_samples_s
 District: Akola, Target Column: Precipitation (mm), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'min_samples_spli
 District: Akola, Target Column: Evapotranspiration (mm), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_samples
 District: Kolhapur, Target Column: Area, Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100, 'min_samples_split': 2, 'min
 District: Kolhapur, Target Column: Production, Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'min_samples_split': 2
 District: Kolhapur, Target Column: Yield, Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_samples_split': 5, 'mi
 District: Kolhapur, Target Column: Irrigated Area, Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100, 'min_samples_split
 District: Kolhapur, Target Column: Annual Rainfall, Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 300, 'min_samples_spli
 District: Kolhapur, Target Column: NITROGEN CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'min_sa
 District: Kolhapur, Target Column: NITROGEN SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'mi
 District: Kolhapur, Target Column: NITROGEN PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200
 District: Kolhapur, Target Column: NITROGEN PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100
 District: Kolhapur, Target Column: PHOSPHATE CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min
 District: Kolhapur, Target Column: PHOSPHATE SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200,
 District: Kolhapur, Target Column: PHOSPHATE PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 20
 District: Kolhapur, Target Column: PHOSPHATE PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 10
 District: Kolhapur, Target Column: POTASH CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100, 'min_sa
 District: Kolhapur, Target Column: POTASH SHARE IN NPK (Percent), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'mi
 District: Kolhapur, Target Column: POTASH PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 100,
 District: Kolhapur, Target Column: POTASH PER HA OF GCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300,
 District: Kolhapur, Target Column: TOTAL CONSUMPTION (tons), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 300, 'min_sa
 District: Kolhapur, Target Column: TOTAL PER HA OF NCA (Kg per ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'mi
 District: Kolhapur, Target Column: FOREST AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 100, 'min_sample
 District: Kolhapur, Target Column: BARREN AND UNCULTIVABLE LAND AREA (1000 ha), Best Hyperparameters: {'subsample': 0.8, 'n_estim
 District: Kolhapur, Target Column: LAND PUT TO NONAGRICULTURAL USE AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_esi
 District: Kolhapur, Target Column: CULTIVABLE WASTE LAND AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 3
 District: Kolhapur, Target Column: PERMANENT PASTURES AREA (1000 ha), Best Hyperparameters: {'subsample': 1.0, 'n estimators': 200,

```
# Calculate and print aggregate metrics for each district
aggregate_metrics = []

for district in districts:
    district_results = [result for result in fine_tuning_results if result['District'] == district]

    if district_results:
        avg_r2 = np.mean([res['R^2'] for res in district_results])
        avg_mse = np.mean([res['MSE'] for res in district_results])
```

```

count = len(district_results)

aggregate_metrics.append({
    'District': district,
    'Average R22: {avg_r2:.4f}")
print(f" Average MSE: {avg_mse:.4f}")
print(f" Models Tuned: {count}")
print("-" * 40)

# Convert aggregate metrics to DataFrame for further use if needed
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)

```

```

☒ District: Akola
  Average R2: 0.9779
  Average MSE: 2254.9571
  Models Tuned: 33
-----
District: Kolhapur
  Average R2: 0.9556
  Average MSE: 173.5197
  Models Tuned: 33
-----
District: Pune
  Average R2: 0.9548
  Average MSE: 4612.1432
  Models Tuned: 33
-----
District: Ratnagiri
  Average R2: 0.9509
  Average MSE: 136278.4674
  Models Tuned: 33
-----
District: Wardha
  Average R2: 0.9558
  Average MSE: 297.8104
  Models Tuned: 33
-----
```

▼ Plot

```

import matplotlib.pyplot as plt
import seaborn as sns

# Set seaborn style for plots
sns.set(style="whitegrid")

# Iterate over districts and target columns to visualize fitting and forecasting
for district in districts:
    district_data = data[data['Dist Name'] == district]
    forecasted_district_data = forecasted_df[forecasted_df['District'] == district]

    for target_column in target_columns:
        # Prepare the actual data
        if target_column not in district_data.columns:
            continue
        actual_data = district_data[['Year', target_column]].dropna()

```

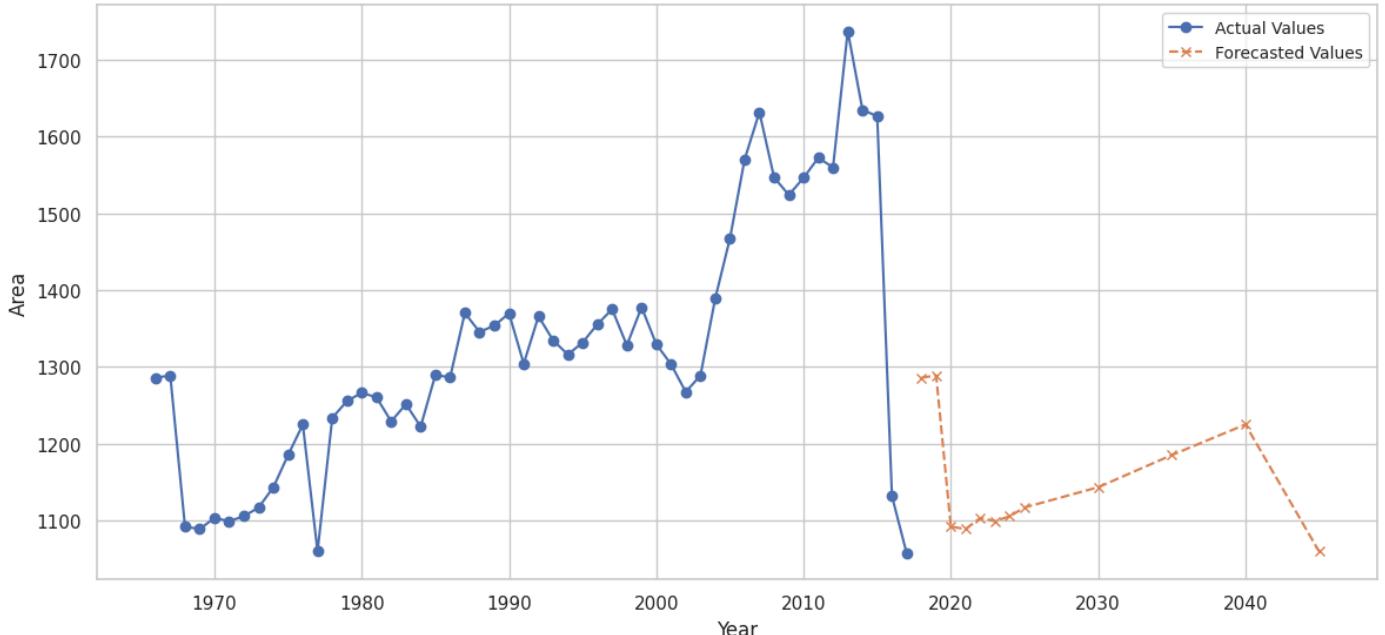
```
# Prepare the forecasted data
forecasted_data = forecasted_district_data[forecasted_district_data['Target Column'] == t]

# Plot the data
plt.figure(figsize=(12, 6))
plt.plot(
    actual_data['Year'], actual_data[target_column], label="Actual Values", marker='o'
)
plt.plot(
    forecasted_data['Year'], forecasted_data['Forecasted Value'], label="Forecasted Value"
)

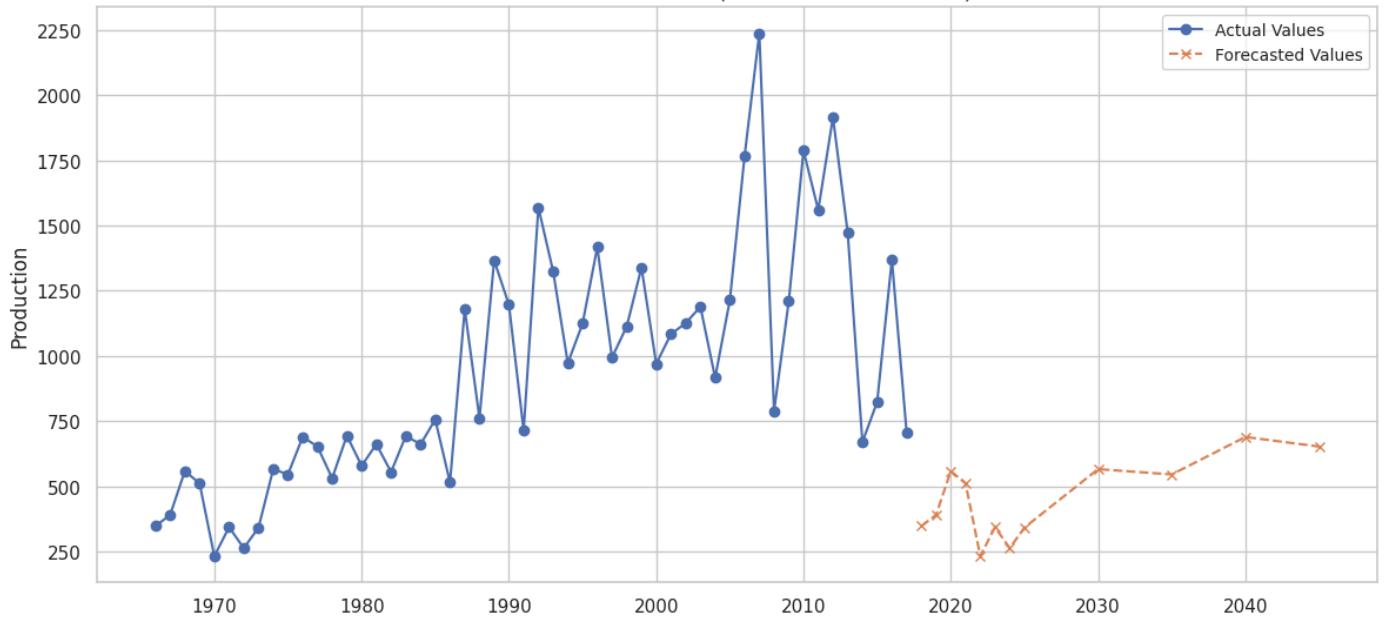
# Add title, labels, and legend
plt.title(f"{district} - {target_column} (Actual vs Forecasted)", fontsize=14)
plt.xlabel("Year", fontsize=12)
plt.ylabel(target_column, fontsize=12)
plt.legend(loc="best", fontsize=10)
plt.grid(True)
plt.tight_layout()
plt.show()
```



Akola - Area (Actual vs Forecasted)



Akola - Production (Actual vs Forecasted)



```
from sklearn.model_selection import train_test_split

# Iterate over districts and target columns to visualize train-test fitting and forecasting
for district in districts:
    district_data = data[data['Dist Name'] == district]

    for target_column in target_columns:
        if target_column not in district_data.columns:
            continue

        # Prepare features and target
        X = district_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')
        y = district_data[target_column]

        if len(y) < 5: # Skip if insufficient data
            continue
```

```

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the Gradient Boosting model
gbr = GradientBoostingRegressor(random_state=42)
gbr.fit(X_train, y_train)

# Predictions on train and test sets
train_predictions = gbr.predict(X_train)
test_predictions = gbr.predict(X_test)

# Plot actual vs predicted for train and test data
plt.figure(figsize=(14, 6))

# Train data
plt.scatter(y_train, train_predictions, label='Train Data', color='blue', alpha=0.6)
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], color='black', lines

# Test data
plt.scatter(y_test, test_predictions, label='Test Data', color='orange', alpha=0.6)

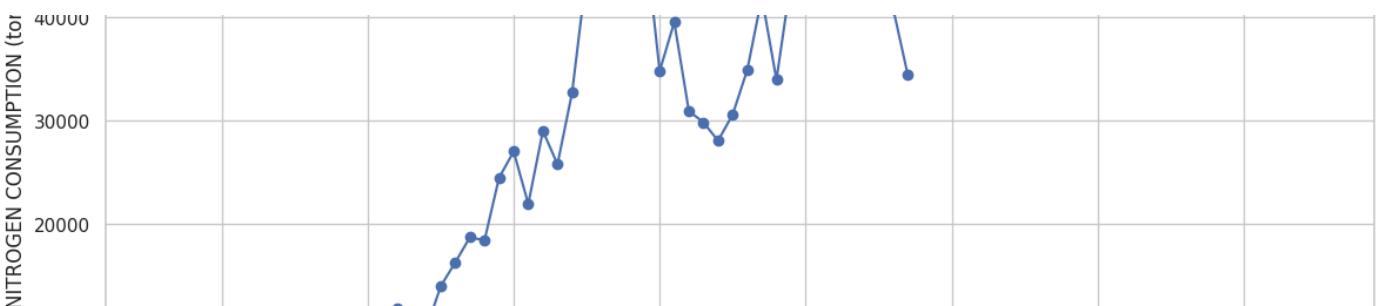
# Add title, labels, and legend
plt.title(f"{district} - {target_column} (Train vs Test Predictions)", fontsize=14)
plt.xlabel("Actual Values", fontsize=12)
plt.ylabel("Predicted Values", fontsize=12)
plt.legend(loc="best", fontsize=10)
plt.grid(True)
plt.tight_layout()
plt.show()

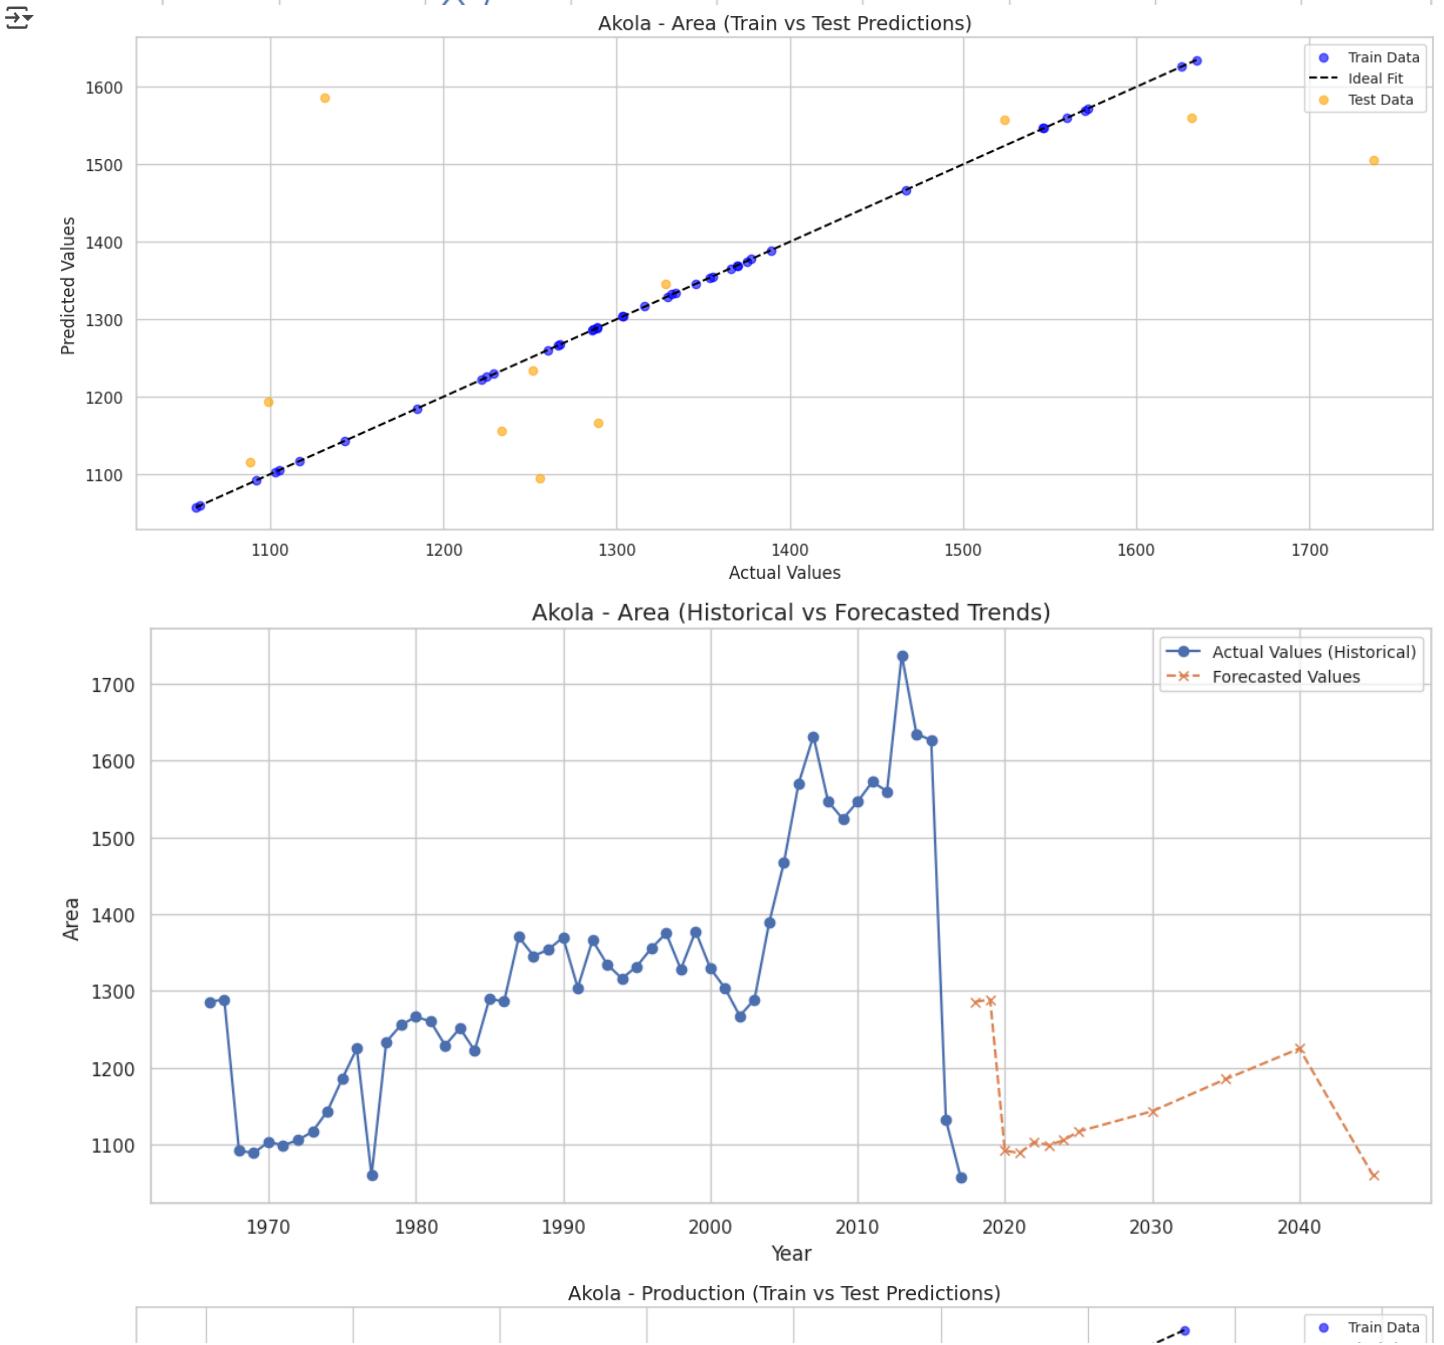
# Forecasting future trends
forecasted_district_data = forecasted_df[forecasted_df['District'] == district]
forecasted_data = forecasted_district_data[forecasted_district_data['Target Column'] == t
actual_data = district_data[['Year', target_column]].dropna()

# Plot historical and forecasted trends
plt.figure(figsize=(12, 6))
plt.plot(actual_data['Year'], actual_data[target_column], label="Actual Values (Historical")
plt.plot(forecasted_data['Year'], forecasted_data['Forecasted Value'], label="Forecasted

# Add title, labels, and legend
plt.title(f"{district} - {target_column} (Historical vs Forecasted Trends)", fontsize=14)
plt.xlabel("Year", fontsize=12)
plt.ylabel(target_column, fontsize=12)
plt.legend(loc="best", fontsize=10)
plt.grid(True)
plt.tight_layout()
plt.show()

```





```

import matplotlib.pyplot as plt
import seaborn as sns

# Set seaborn style for the plots
sns.set(style="whitegrid")

# Get unique districts
districts = forecasted_df['District'].unique()

# Iterate over districts to create individual plots
for district in districts:
    # Filter data for the specific district
    district_data = forecasted_df[forecasted_df['District'] == district]

    # Create a new plot for the district
    plt.figure(figsize=(12, 6))

```

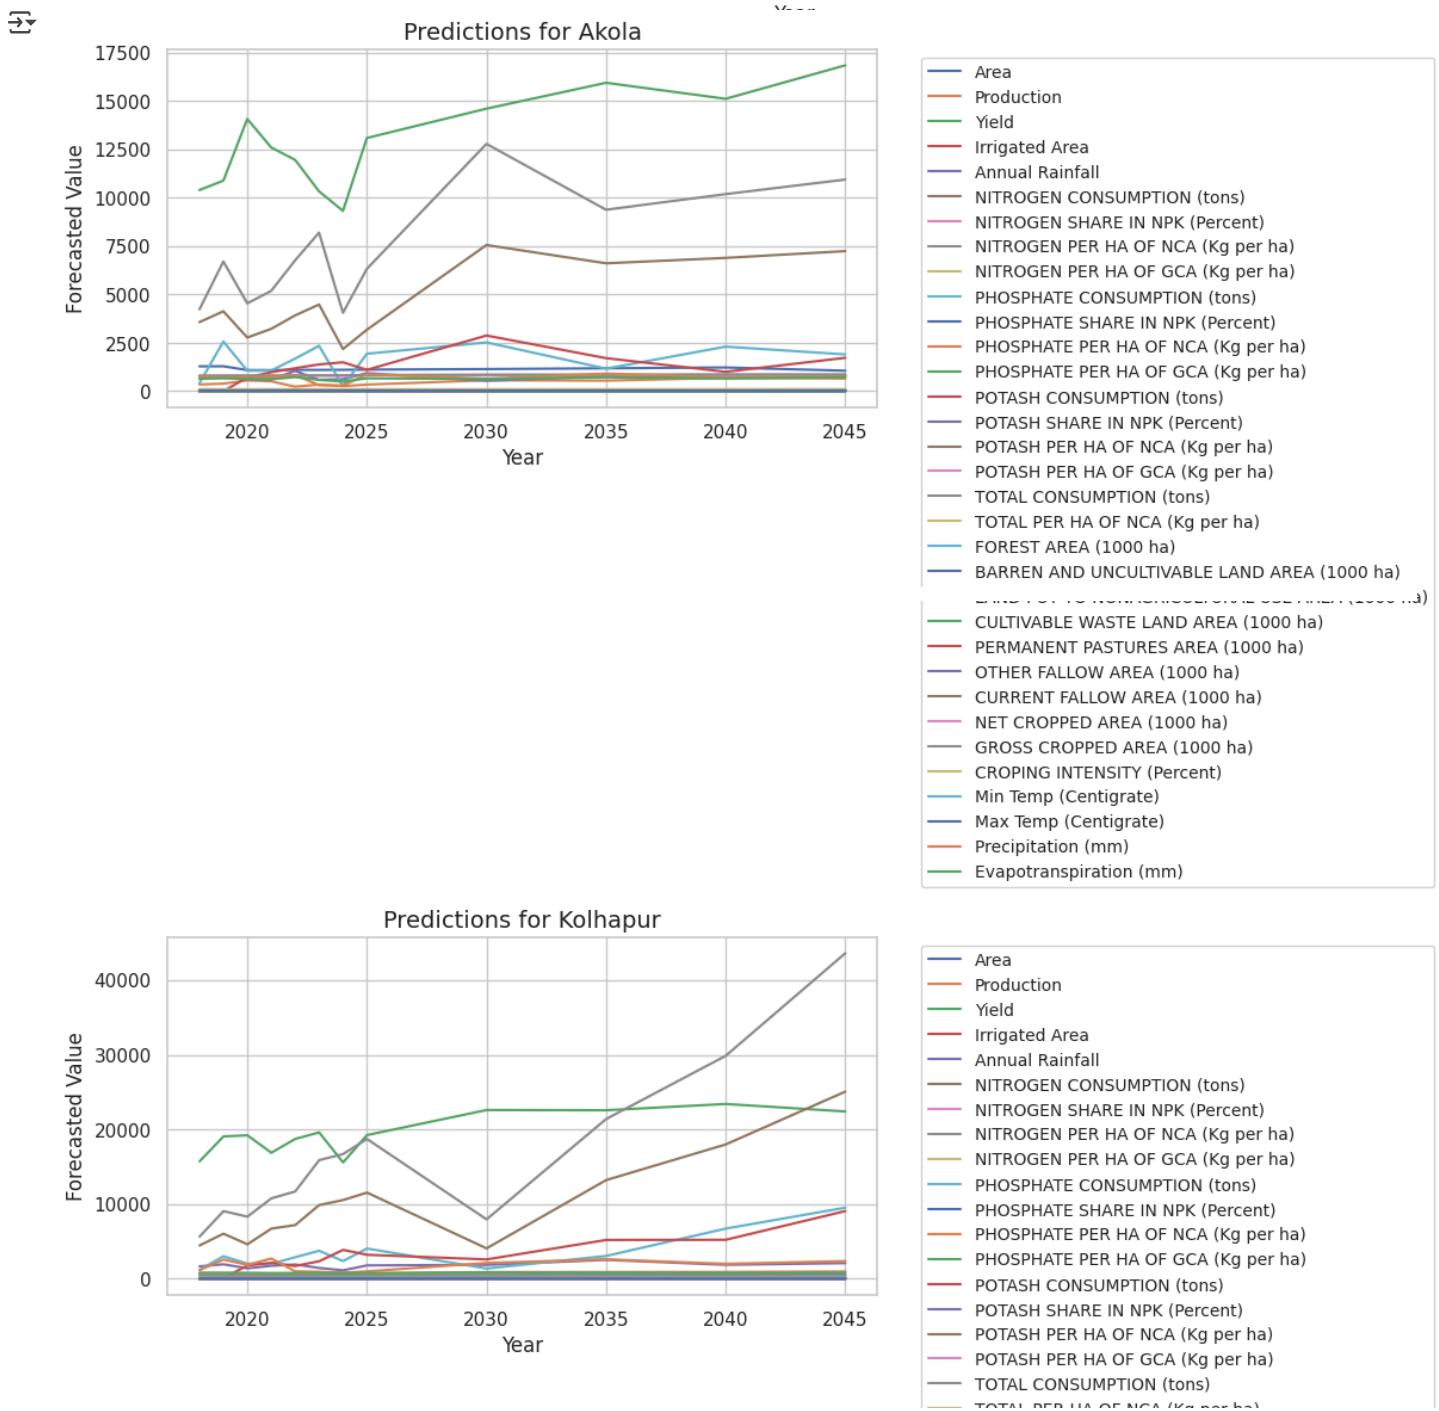
```

# Plot each target column
target_columns = district_data['Target Column'].unique()
for target in target_columns:
    target_data = district_data[district_data['Target Column'] == target]
    plt.plot(target_data['Year'], target_data['Forecasted Value'], label=target)

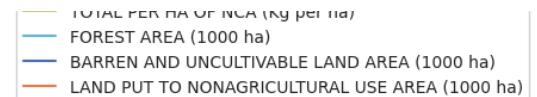
# Add title, labels, and legend
plt.title(f"Predictions for {district}", fontsize=14)
plt.xlabel("Year", fontsize=12)
plt.ylabel("Forecasted Value", fontsize=12)
plt.legend(loc="upper left", bbox_to_anchor=(1.05, 1), fontsize=10) # Legend outside the plot

# Show the plot
plt.tight_layout()
plt.show()

```



✓ 18 Jan



```
import gdown
file_id = '1KXMl4EVNofFnGsypt1UpZ4GJL6pGLJCe'
url = f'https://drive.google.com/uc?id={file_id}'
output = 'maharashtra_aggregate_1966_2017.csv'
gdown.download(url, output, quiet=False)
import pandas as pd
df = pd.read_csv('maharashtra_aggregate_1966_2017.csv')
```

⬇️ Downloading...
From: <https://drive.google.com/uc?id=1KXMl4EVNofFnGsypt1UpZ4GJL6pGLJCe>

To: /content/maharashtra_aggregate_1966_2017.csv
100% [██████████] 277k/277k [00:00<00:00, 74.6MB/s]

LARGE SCALE INDUSTRY (LARGE)



Start coding or generate with AI.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

Load the dataset

```
data = pd.read_csv('maharashtra_aggregate_1966_2017.csv')
data = data.drop(columns=['TOTAL AREA (1000 ha)', 'State Name'])
```

List of districts to forecast
districts = [

```
    "Ahmednagar", "Akola", "Amarawati", "Aurangabad", "Beed", "Bhandara",
    "Buldhana", "Chandrapur", "Dhule", "Jalgaon", "Kolhapur", "Nagpur",
    "Nanded", "Nasik", "Osmanabad", "Parbhani", "Pune", "Sangli",
    "Satara", "Solapur", "Yeotmal"
]
```

target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude 'Year' and 'Dist Name'

Define a function to calculate additional metrics
def calculate_metrics(y_true, y_pred):

```
    r2 = r2_score(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_true, y_pred)
    mape = np.mean(np.abs((y_true - y_pred) / y_true)) * 100
    return {'R2': r2, 'MSE': mse, 'RMSE': rmse, 'MAE': mae, 'MAPE': mape}
```

Dictionary to store results
results = {}

```
# Iterate over each district
# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]
    results[district] = {}

# Loop through each target column
```

```

for target_column in target_columns:
    print(f"  Forecasting: {target_column}")

    # Prepare features (X) and target (y)
    X = district_data.drop(columns=['Year', 'Dist Name', target_column])
    y = district_data[target_column]

    # Check if X and y are empty
    if X.empty or y.empty:
        print(f"    Skipping {target_column} for {district} due to insufficient data.")
        continue # Skip to the next target column

    # Train-test split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Gradient Boosting Model
    model = GradientBoostingRegressor(random_state=42)
    model.fit(X_train, y_train)

    # Predictions
    y_pred = model.predict(X_test)

    # Evaluate the model
    metrics = calculate_metrics(y_test, y_pred)
    results[district][target_column] = metrics

# Print results
for district, district_metrics in results.items():
    print(f"\nResults for {district}:")
    for column, metrics in district_metrics.items():
        print(f"  {column}:")
        for metric, value in metrics.items():
            print(f"    {metric}: {value:.4f}")

# Dictionary to store all forecasted values
forecasted_values = []

# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]
    results[district] = {}

    # Loop through each target column
    for target_column in target_columns:
        print(f"  Forecasting: {target_column}")

        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column])
        y = district_data[target_column]

        # Check if there is sufficient data for this target column
        if X.empty or y.empty:
            print(f"    Skipping {target_column} for {district} due to insufficient data.")
            continue # Skip to the next target column

        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

# Gradient Boosting Model
model = GradientBoostingRegressor(random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
metrics = calculate_metrics(y_test, y_pred)
results[district][target_column] = metrics

# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    district_data = data[data['Dist Name'] == district]
    results[district] = {}

    # Loop through each target column
    for target_column in target_columns:
        print(f"  Forecasting: {target_column}")

        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column])
        y = district_data[target_column]

        # Check if there is sufficient data for this target column
        if X.empty or y.empty:
            print(f"    Skipping {target_column} for {district} due to insufficient data.")
            continue # Skip to the next target column

        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Store results
        for i in range(len(X_test)):
            forecasted_values.append({
                'District': district,
                'Target Column': target_column,
                'Actual Value': y_test.iloc[i],
                'Predicted Value': y_pred[i]
            })

    # Convert forecasted values to DataFrame
    forecasted_df = pd.DataFrame(forecasted_values)

    # Save to CSV
    forecasted_csv_path = '/content/forecasted_values_18Jan.csv'
    forecasted_df.to_csv(forecasted_csv_path, index=False)

print(f"Forecasted values saved to {forecasted_csv_path}")

# Prepare data for performance metrics
metrics_list = []

for district, district_metrics in results.items():
    for column, metrics in district_metrics.items():

```

```

metrics_row = {'District': district, 'Target Column': column}
metrics_row.update(metrics) # Add all metrics
metrics_list.append(metrics_row)

# Convert metrics list to DataFrame
metrics_df = pd.DataFrame(metrics_list)

# Save to CSV
metrics_csv_path = '/content/performance_metrics_18Jan.csv'
metrics_df.to_csv(metrics_csv_path, index=False)

print(f"Performance metrics saved to {metrics_csv_path}")

# Initialize a list to store aggregate metrics
aggregate_metrics = []

# Iterate through each district
# Initialize a list to store aggregate metrics
aggregate_metrics = []

# Iterate through each district
for district in results.keys():
    district_metrics = results[district]

    # Initialize accumulators for each metric
    r2_sum = 0
    mse_sum = 0
    rmse_sum = 0
    mae_sum = 0
    mape_sum = 0
    count = 0

    # Aggregate metrics across all target columns
    for target_column, metrics in district_metrics.items():
        r2_sum += metrics['R2']
        mse_sum += metrics['MSE']
        rmse_sum += metrics['RMSE']
        mae_sum += metrics['MAE']
        mape_sum += metrics['MAPE']
        count += 1 # Increment count inside the target column loop

    # Calculate averages, but only if count is greater than 0
    if count > 0: # Check to avoid division by zero
        aggregate_metrics.append({
            'District': district,
            'Avg R2

```

```

aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)

print(f"Aggregate performance metrics saved to {aggregate_metrics_csv_path}")

# Display the DataFrame
aggregate_metrics_df

import pandas as pd
import numpy as np
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split

# Load the dataset
data = pd.read_csv('maharashtra_aggregate_1966_2017.csv')
data = data.drop(columns=['TOTAL AREA (1000 ha)', 'State Name'])
# Years to forecast
future_years = list(range(2018, 2026)) + [2030, 2035, 2040, 2045]

# List of districts to forecast
districts = [
    "Ahmednagar", "Akola", "Amarawati", "Aurangabad", "Beed", "Bhandara",
    "Buldhana", "Chandrapur", "Dhule", "Jalgaon", "Kolhapur", "Nagpur",
    "Nanded", "Nasik", "Osmanabad", "Parbhani", "Pune", "Sangli",
    "Satara", "Solapur", "Yeotmal"
]
# Replace with actual district names
target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude 'Year' and 'Dist Name'

# Dictionary to store forecasted results
forecasted_results = []

# Iterate over each district
for district in districts:
    print(f"Processing District: {district}")
    # Ensure district name is consistent with the DataFrame
    district_data = data[data['Dist Name'].str.strip() == district.strip()]

    # Check if district_data is empty before proceeding
    if district_data.empty:
        print(f"Skipping {district} due to no matching data.")
        continue

    results[district] = {}

    for target_column in target_columns:
        print(f"  Forecasting: {target_column}")

        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column])
        y = district_data[target_column]

        # Train-test split (train on all historical data for forecasting)
        X_train = X
        y_train = y

        # Gradient Boosting Model
        model = GradientBoostingRegressor(random_state=42)

```

```

model.fit(X_train, y_train)

# Create future data for predictions
future_data = district_data.iloc[:len(future_years)].copy()
future_data['Year'] = future_years
future_X = future_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')

# Predict future values
future_predictions = model.predict(future_X)

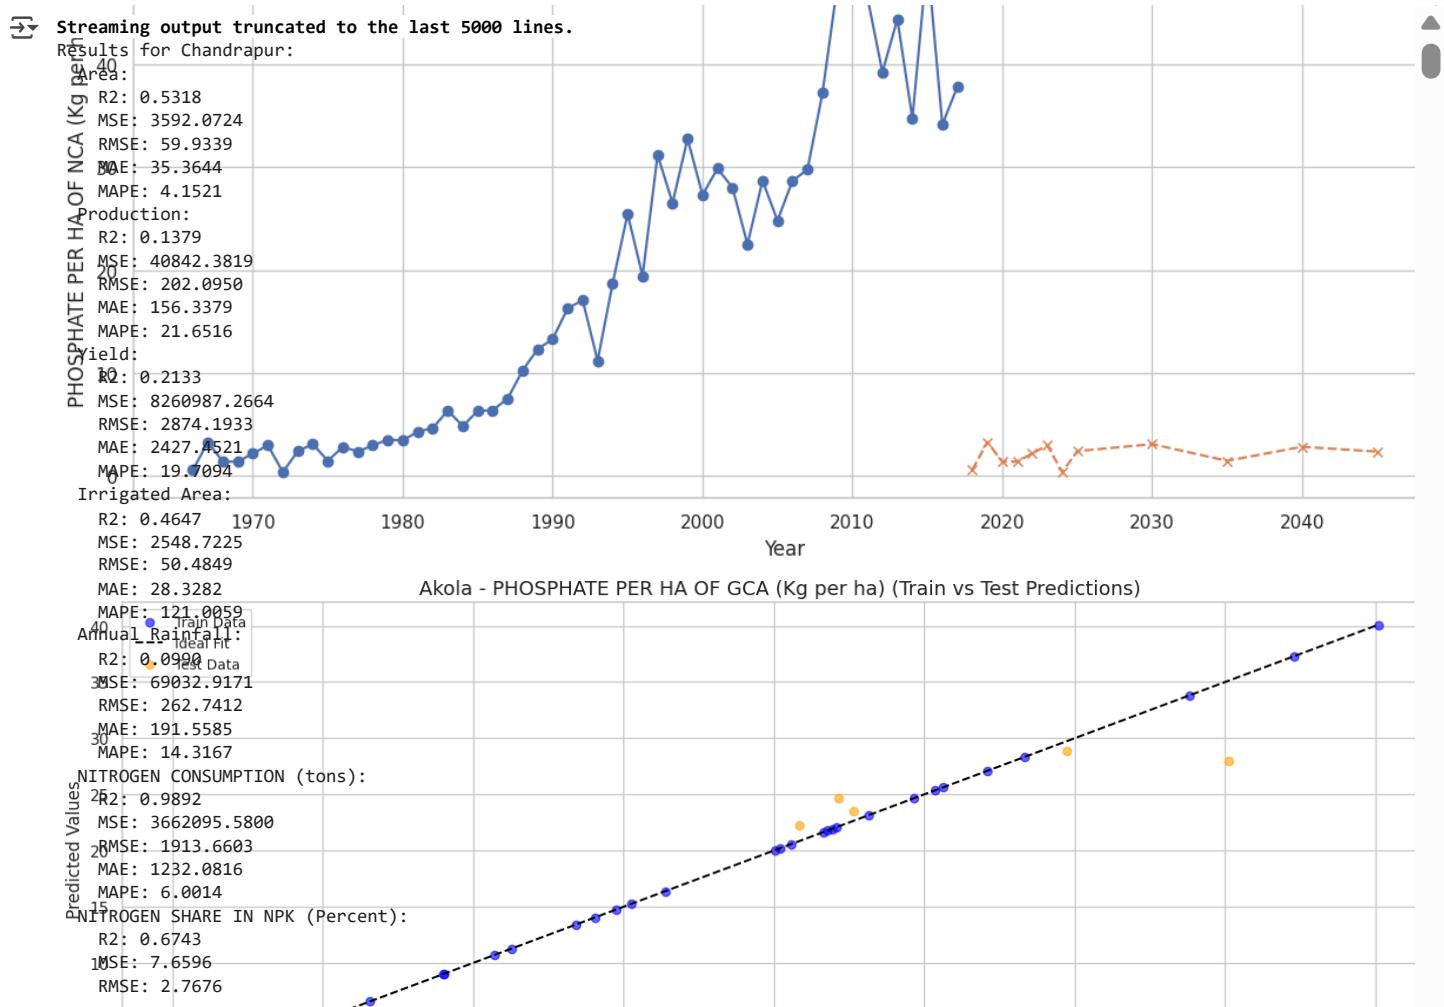
# Store results
for year, prediction in zip(future_years, future_predictions):
    forecasted_results.append({
        'District': district,
        'Year': year,
        'Target Column': target_column,
        'Forecasted Value': prediction
    })

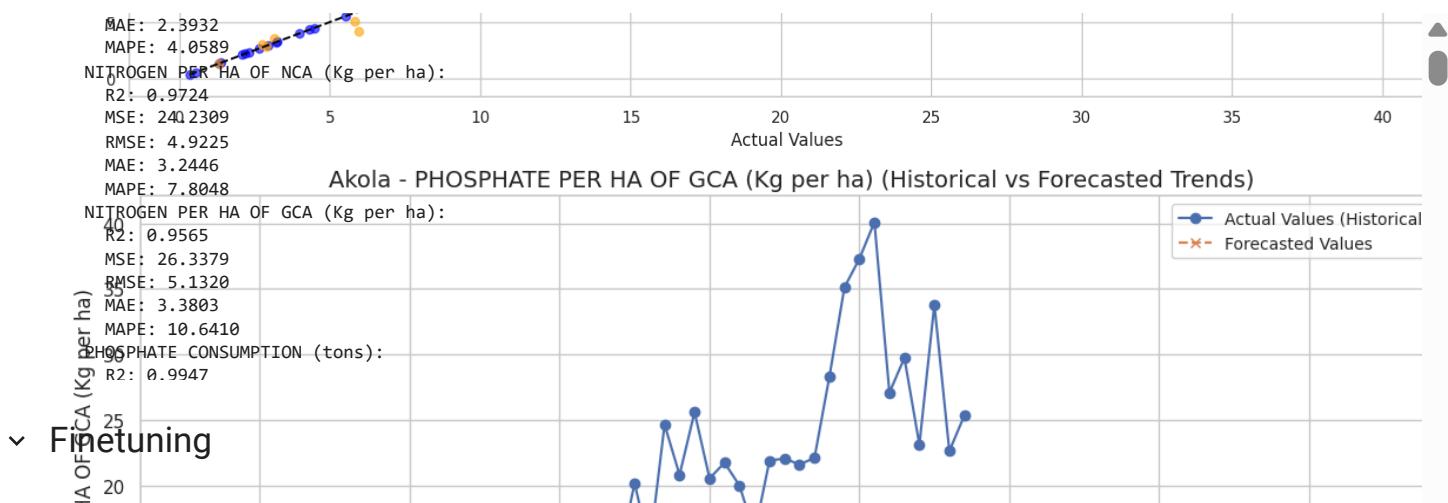
# Convert forecasted results to DataFrame
forecasted_df = pd.DataFrame(forecasted_results)

# Save to CSV
forecasted_csv_path = '/content/18Jan_forecasted_future_values.csv'
forecasted_df.to_csv(forecasted_csv_path, index=False)

print(f"Forecasted values saved to {forecasted_csv_path}")

```





```

from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd
import numpy as np
from joblib import Parallel, delayed

# Load the dataset
file_path = '/content/maharashtra_aggregate_1966_2017.csv'
data = pd.read_csv(file_path)
data = data.drop(columns=['TOTAL AREA (1000 ha)', 'State Name']) # Adjust as per actual column name

# Future years to forecast
future_years = list(range(2018, 2026)) + [2030, 2035, 2040, 2045]

# Districts to process
districts = [
    "Ahmednagar", "Akola", "Amarawati", "Aurangabad", "Beed", "Bhandara",
    "Buldhana", "Chandrapur", "Dhule", "Jalgaon", "Kolhapur", "Nagpur",
    "Nanded", "Nasik", "Osmanabad", "Parbhani", "Pune", "Sangli",
    "Satara", "Solapur", "Yeotmal"
]

target_columns = data.columns.drop(['Year', 'Dist Name']) # Exclude non-target columns

# Hyperparameter tuning grid
param_distributions = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 5],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 3],
    'subsample': [0.8, 1.0]
}

# Function to process each district and target column
def process_district(district):
    results = {'fine_tuning': [], 'forecasting': []}
    district_data = data[data['Dist Name'] == district]

    for target_column in target_columns:
        # Prepare features (X) and target (y)
        X = district_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')

```

```

y = district_data[target_column]

# Skip columns with insufficient data
if len(y) < 5:
    continue

# Train on historical data
X_train = X
y_train = y

# Randomized search for hyperparameter tuning
gbr = GradientBoostingRegressor(random_state=42)
random_search = RandomizedSearchCV(
    estimator=gbr,
    param_distributions=param_distributions,
    n_iter=10, # Reducing iterations for speed
    cv=2,
    scoring='neg_mean_squared_error',
    n_jobs=-1,
    verbose=0
)
random_search.fit(X_train, y_train)

# Best parameters and model
best_params = random_search.best_params_
best_model = random_search.best_estimator_

# Forecasting future years
future_data = district_data.iloc[:len(future_years)].copy()
future_data['Year'] = future_years
future_X = future_data.drop(columns=['Year', 'Dist Name', target_column], errors='ignore')
future_predictions = best_model.predict(future_X)

# Store fine-tuning results
mse = mean_squared_error(y_train, best_model.predict(X_train))
r2 = r2_score(y_train, best_model.predict(X_train))
results['fine_tuning'].append({
    'District': district,
    'Target Column': target_column,
    'Best Parameters': best_params,
    'MSE': mse,
    'R2

```

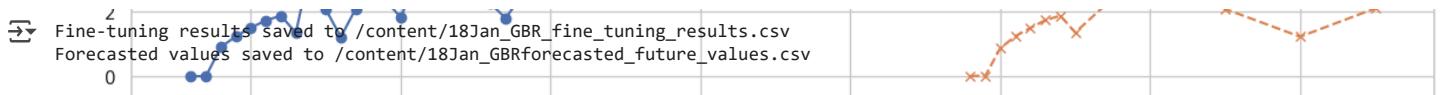
```

# Combine results
fine_tuning_results = [item for result in all_results for item in result['fine_tuning']]
forecasted_results = [item for result in all_results for item in result['forecasting']]

# Save fine-tuning metrics to CSV
metrics_df = pd.DataFrame(fine_tuning_results)
metrics_csv_path = '/content/18Jan_GBR_fine_tuning_results.csv'
metrics_df.to_csv(metrics_csv_path, index=False)
print(f"Fine-tuning results saved to {metrics_csv_path}")

# Save forecasted results to CSV
forecasted_df = pd.DataFrame(forecasted_results)
forecasted_csv_path = '/content/18Jan_GBRforecasted_future_values.csv'
forecasted_df.to_csv(forecasted_csv_path, index=False)
print(f"Forecasted values saved to {forecasted_csv_path}")

```



```

# Initialize lists for aggregate metrics and forecast summaries
aggregate_metrics = []
forecast_summaries = []

# Aggregate fine-tuning results by district
for district in metrics_df['District'].unique():
    district_metrics = metrics_df[metrics_df['District'] == district]

    # Calculate average MSE and R2 for the district
    avg_mse = district_metrics['MSE'].mean()
    avg_r2 = district_metrics['R2'].mean()

    # Append aggregate metrics for the district
    aggregate_metrics.append({
        'District': district,
        'Avg MSE': avg_mse,
        'Avg R2

```

```

forecast_summary_df = pd.DataFrame(forecast_summaries)

# Save aggregated metrics to CSV
aggregate_metrics_csv_path = '/content/18Jan_GBRaggregate_fine_tuning_metrics.csv'
aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)
print(f"Aggregate fine-tuning metrics saved to {aggregate_metrics_csv_path}")

# Save forecast summaries to CSV
forecast_summary_csv_path = '/content/18Jan_GBRforecast_summary.csv'
forecast_summary_df.to_csv(forecast_summary_csv_path, index=False)
print(f"Forecast summaries saved to {forecast_summary_csv_path}")

# Display the DataFrames
print("\nAggregate Fine-Tuning Metrics:")
print(aggregate_metrics_df)

print("\nForecast Summary:")
print(forecast_summary_df)

# Initialize a list to store aggregate metrics
aggregate_metrics = []

# Get unique districts from the results
# Changed 'results_df' to 'metrics_df'
unique_districts = metrics_df['District'].unique()

# Iterate through each district
for district in unique_districts:
    # Filter results for the current district
    # Changed 'results_df' to 'metrics_df'
    district_data = metrics_df[metrics_df['District'] == district]

    # Calculate average metrics for the district
    avg_rmse = district_data['MSE'].mean() # Assuming 'MSE' is the desired metric
    avg_r2 = district_data['R^2'].mean()

    # Append aggregate metrics for the district
    aggregate_metrics.append({
        'District': district,
        'Avg RMSE': avg_rmse, # Or any other relevant metric
        'Avg R^2': avg_r2
    })

# Convert aggregate metrics to a DataFrame
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)

# Save to CSV if needed
aggregate_metrics_csv_path = '/content/18Jan_GBRaggregate_performance_metrics.csv'
aggregate_metrics_df.to_csv(aggregate_metrics_csv_path, index=False)

print(f"Aggregate performance metrics saved to {aggregate_metrics_csv_path}")

# Display the DataFrame
aggregate_metrics_df

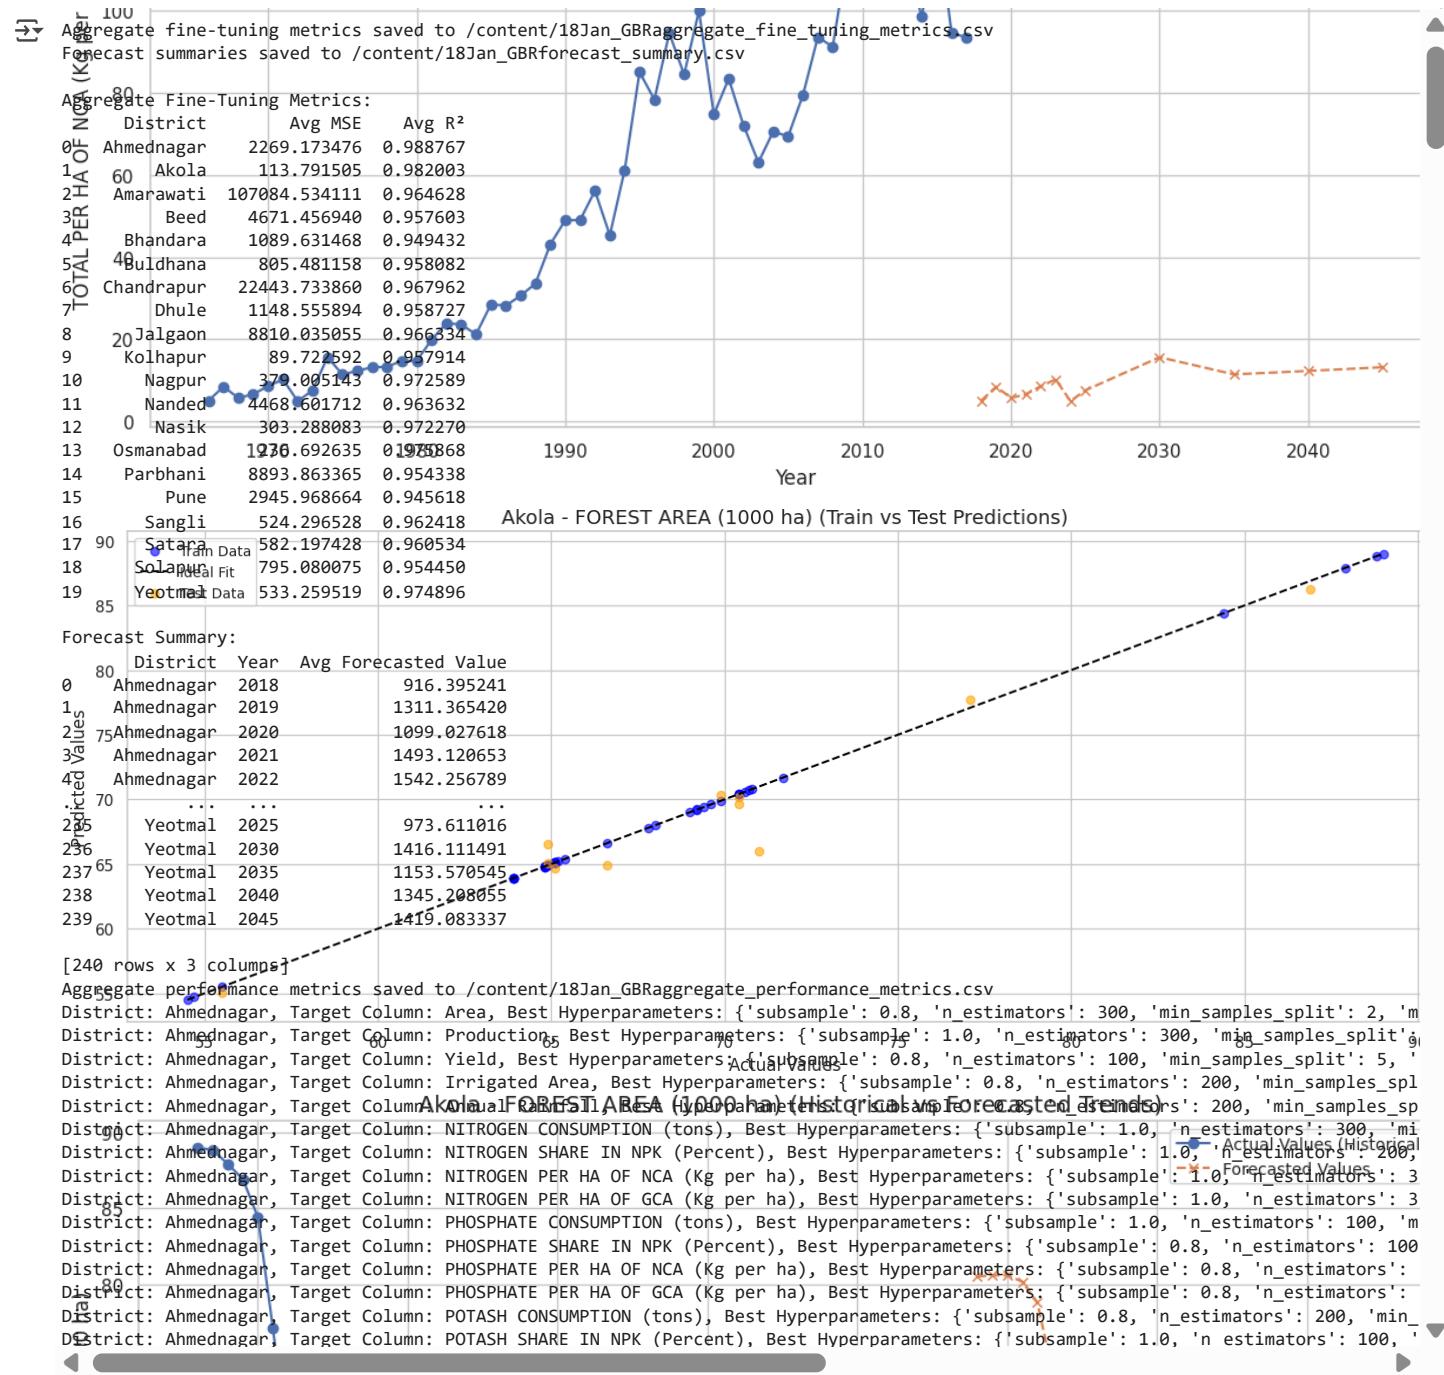
# Print the best hyperparameters for each district and target column
for result in fine_tuning_results:
    district = result['District']

```

```

target_column = result['Target Column']
best_params = result['Best Parameters']
print(f'District: {district}, Target Column: {target_column}, Best Hyperparameters: {best_params}')

```



```

# Calculate and print aggregate metrics for each district
aggregate_metrics = []

```

```

for district in districts:
    district_results = [result for result in fine_tuning_results if result['District'] == district]

    if district_results:
        avg_r2 = np.mean([res['R2'] for res in district_results])
        avg_mse = np.mean([res['MSE'] for res in district_results])
        count = len(district_results)

        aggregate_metrics.append({
            'District': district,
            'Avg R2': avg_r2,
            'Avg MSE': avg_mse
        })

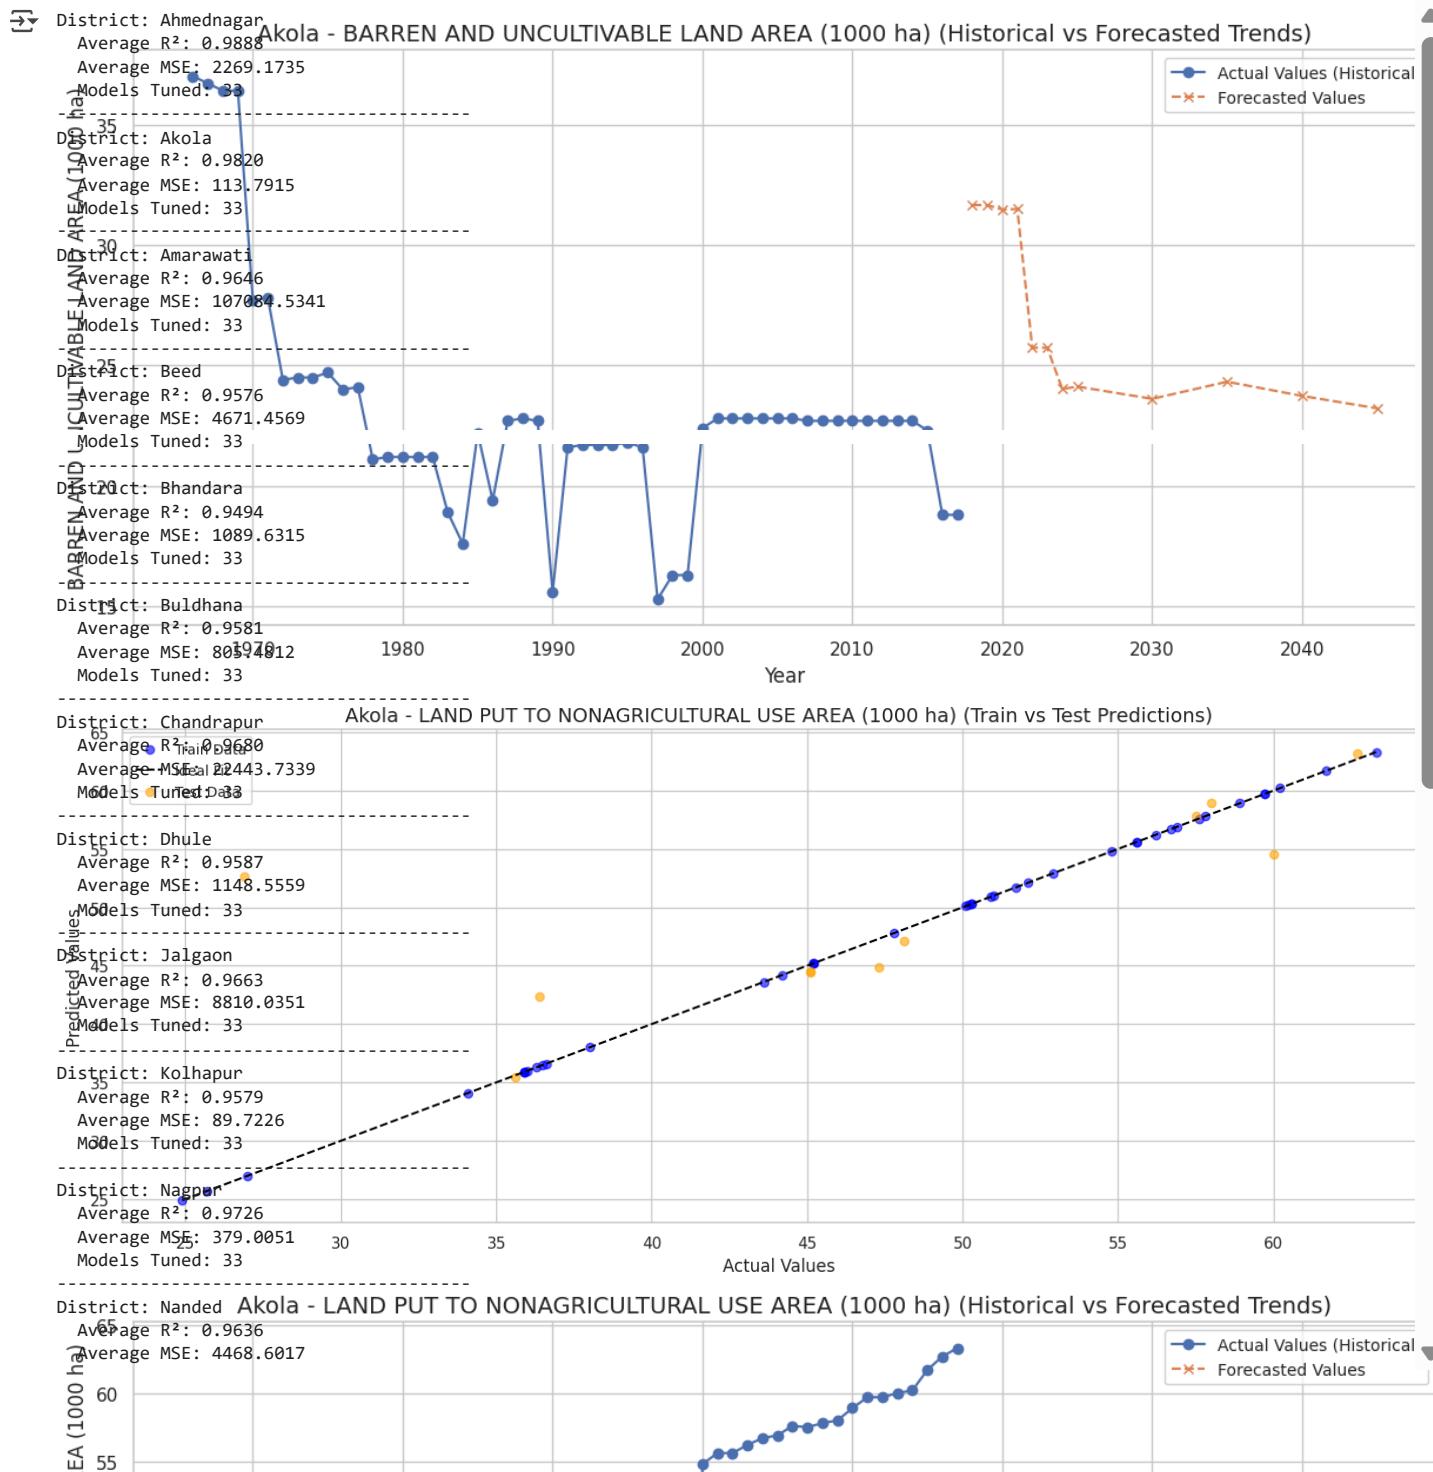
```

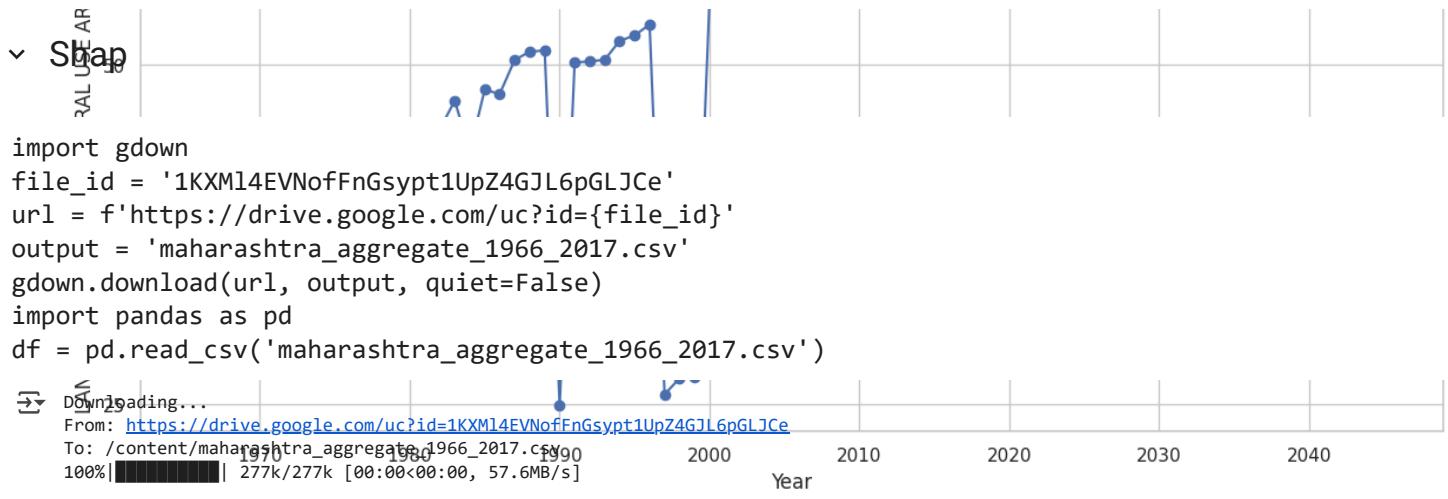
```

        'Average R22: {avg_r2:.4f}")
    print(f"  Average MSE: {avg_mse:.4f}")
    print(f"  Models Tuned: {count}")
    print("-" * 40)

# Convert aggregate metrics to DataFrame for further use if needed
aggregate_metrics_df = pd.DataFrame(aggregate_metrics)

```





Start coding or generate with AI.

```
# https://drive.google.com/file/d/1kVWW79nASA1FImE5FhEXjdy4X_jIPHJF/view?usp=drive_link
import gdwn
file_id = '1kVWW79nASA1FImE5FhEXjdy4X_jIPHJF'
url = f'https://drive.google.com/uc?id={file_id}'
output = 'GBR18Jan_forecasted_future_values.csv'
gdwn.download(url, output, quiet=False)
import pandas as pd
df = pd.read_csv('GBR18Jan_forecasted_future_values.csv')
```



```
import pandas as pd
import shap
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
```

```
# Load dataset
data_path = "GBR18Jan_forecasted_future_values.csv"
df = pd.read_csv(data_path)

# Define target and features
targets = ['Yield', 'Production'] # Specify your target columns
features = [col for col in df.columns if col not in targets] # Exclude target columns

# Split data
X = df[features]
y_yield = df['Yield']
y_production = df['Production']

X_train, X_test, y_train_yield, y_test_yield = train_test_split(X, y_yield, test_size=0.2, random_state=42)
X_train, X_test, y_train_production, y_test_production = train_test_split(X, y_production, test_size=0.2, random_state=42)

# Train RandomForestRegressor for Yield
model_yield = RandomForestRegressor(n_estimators=100, random_state=42)
model_yield.fit(X_train, y_train_yield)
```

```

# Train RandomForestRegressor for Production
model_production = RandomForestRegressor(n_estimators=100, random_state=42)
model_production.fit(X_train, y_train_production)

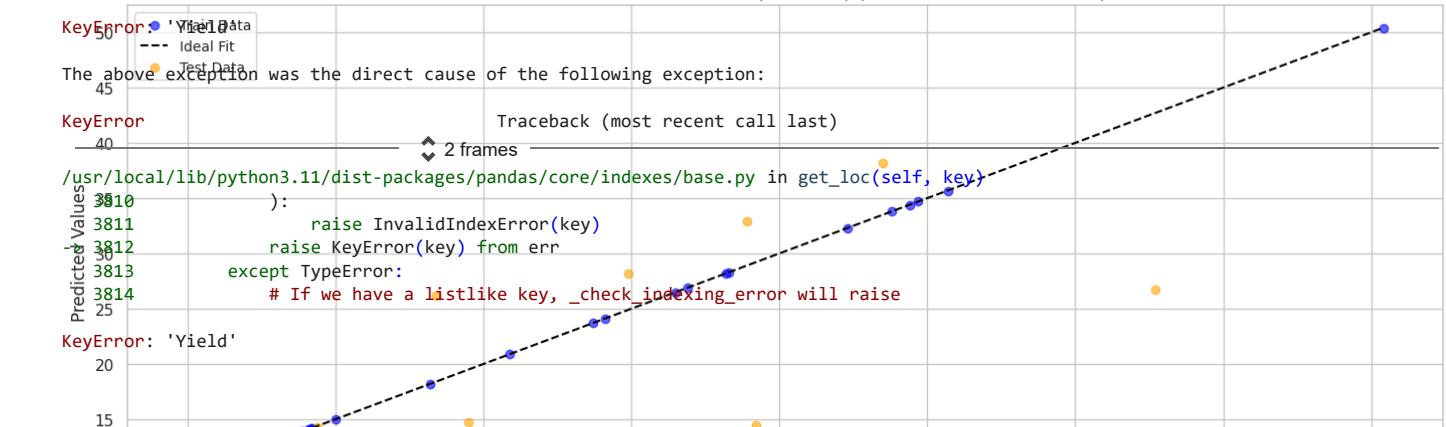
# Apply SHAP for Yield
explainer_yield = shap.Explainer(model_yield, X_train)
shap_values_yield = explainer_yield(X_test)

# Apply SHAP for Production
explainer_production = shap.Explainer(model_production, X_train)
shap_values_production = explainer_production(X_test)

# Plot SHAP summary for Yield
shap.summary_plot(shap_values_yield, X_test, plot_type="bar", show=False)
plt.title("Feature Importance for Yield")
plt.show()

# Plot SHAP summary for Production
shap.summary_plot(shap_values_production, X_test, plot_type="bar", show=False)
plt.title("Feature Importance for Production")
plt.show()

```



```

import pandas as pd
import shap
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor

# Load dataset
data_path = "GBR18Jan_forecasted_future_values.csv"
df = pd.read_csv(data_path)

# Print column names to verify their exact names

```

```
print("Dataset Columns:", df.columns)

# Define target and features
actual_targets = [col for col in df.columns if 'Yield' in col or 'Production' in col] # Find clo
targets = actual_targets if actual_targets else ['Yield', 'Production'] # Default if not found

if not set(targets).issubset(df.columns):
    raise KeyError(f"Columns {targets} not found in dataset. Available columns: {df.columns}")
```